

1 Introduction

The experiment aims at realizing handwritten digit (from 0 to 9) recognition using specific method.

For each handwritten digit, a 32*32 normalized bitmap is generated to represent the digit. Then, the bitmap is divided into blocks. Each block has a size of 4*4 bit, so there are 8*8 blocks in a bitmap for a digit. The number of pixels (ranging from 0 to 16) are counted in each block. Therefore, the input contains 64 numeric attributes.

2 Visualization

To better visualize the generated images without rotation or mirror, I modified the line 75 of the Python script, switching “row” and “col”. That is “digitImage.setPixel(col,row,px)”.

The Python program opens the test file, “optdigits.tes.csv”, as input file. When starting running the program, it firstly prints out the header line of the csv file, as is shown below.

```
x11,x12,x13,x14,x15,x16,x17,x18,x21,x22,x23,x24,x25,x26,x27,x28,x31,x32,x33,x34,
x35,x36,x37,x38,x41,x42,x43,x44,x45,x46,x47,x48,x51,x52,x53,x54,x55,x56,x57,x58,
x61,x62,x63,x64,x65,x66,x67,x68,x71,x72,x73,x74,x75,x76,x77,x78,x81,x82,x83,x84,
x85,x86,x87,x88,digit
```

For an image with 8*8 blocks, “x23” represents the block located at the **second row from the top** and the **third column from the left**. The last one of the header is the real digit which that image shows, regarded as the class label.

Since the Python program sets “lines = lines[871:874]”, it extracts three instances from the csv file and produce corresponding images. The positions of the three instances in “optdigits.tes.csv” are shown in the following screen shot.

873	0	0	0	11	7	0	0	0	0	0	8	15	7	0	0
874	0	0	2	13	15	1	0	0	0	1	14	13	15	4	0
875	0	0	0	6	14	0	0	0	0	0	4	16	6	0	0

For the first image, the Python program prints number “1” followed by an array with 64 values of the 64 blocks, as is shown below. Each value ranges from 0 to 16, representing the number of pixels of that block. Among them, the first appeared “11” which is the fourth element in the array, means that the color intensity of block located at the first row from the top and the fourth column from the left is 11.

```
1 ['0', '0', '0', '11', '7', '0', '0', '0', '0', '0', '0', '8', '15', '7', '0', '0',
  '0', '0', '0', '13', '8', '0', '0', '0', '0', '0', '16', '14', '8', '1', '0',
  '0', '0', '5', '16', '10', '10', '14', '1', '0', '0', '2', '15', '3', '0', '1',
  '2', '7', '0', '0', '0', '10', '13', '1', '10', '11', '0', '0', '0', '0', '10', '1',
  '16', '15', '5', '0', '6']
```

Then, the program prints out five columns for the 64 blocks. The first column is the row number of a block, ranging from 0 to 7. The second column is the column number of a block, also ranging from 0 to 7. The third column is the index of a block, ranging from 0 to 63. The fourth column is the value of the integer of a block, ranging from 0 to 16. The fifth column is the RGB value of a block, ranging from 0 to 255.

```

0 0 0 0 255
0 1 1 0 255
0 2 2 0 255
0 3 3 11 79
0 4 4 7 143
0 5 5 0 255
0 6 6 0 255
0 7 7 0 255
1 0 8 0 255
1 1 9 0 255
1 2 10 8 127
1 3 11 15 15
1 4 12 7 143
1 5 13 0 255
1 6 14 0 255
1 7 15 0 255
2 0 16 0 255
2 1 17 0 255

```

In order to draw a whole image with pixel RGB values, it is necessary to change the scale from 0-16 to 255-0 for each block. The pixel RGB of white color is (255,255,255), and the pixel RGB of black color is (0,0,0). Since it is a bitmap (black-and-white), for each pixel, the RGB three values should be the same ranging from 255 to 0. Then the program adopts the following equation to do the transformation.

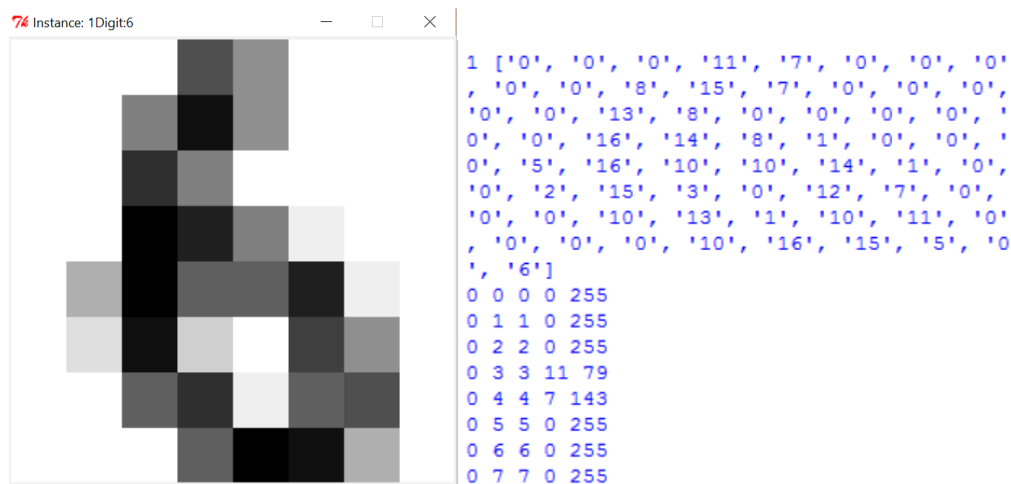
```

#change scale from 0-16 to 255-0
intensity = int(((16-val)/16.0)*255)

```

Then, to better visualize in a larger figure, the Python program blows up the image by a factor of 50.

Finally, to display the image, a window titled with the instance number of the image followed by the correct digit number it represents, is generated, as is shown below.



Compared with the two screen shots, it can be found that the first three blocks on the first row are white colors. The 4th block on the first row has original value “11”, transformed to RGB (79,79,79). The 5th block on the first row has original value “7”, transformed to RGB (143,143,143), a bit lighter than the 4th block on the same row. As a whole, the image indeed looks like a handwritten digit “6”.

After clicking on the first generated image, the image disappeared, and then press enter to view the second image (digit “8”). Then likewise the third image (digit “4”) generates.

3 Replicating the Results

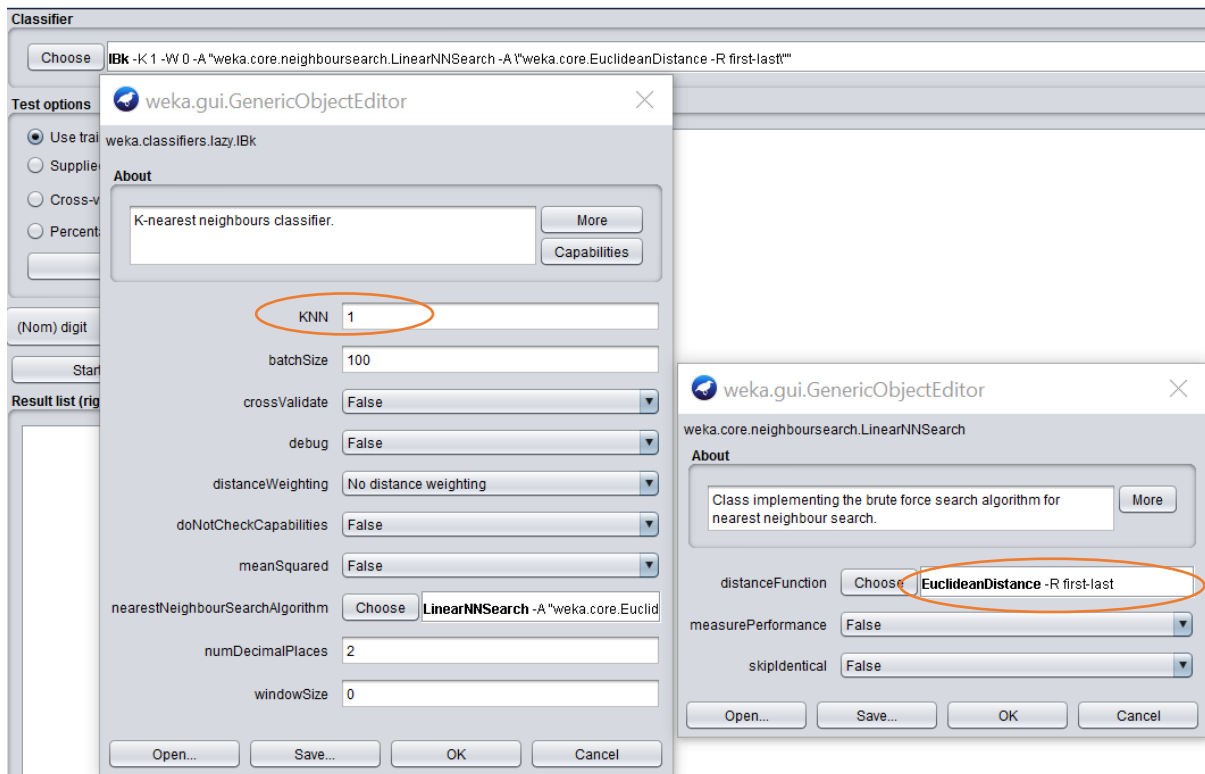
(Weka log is not quite easy to read, so I decide not to append it to my documentation.)

3.1 Methods

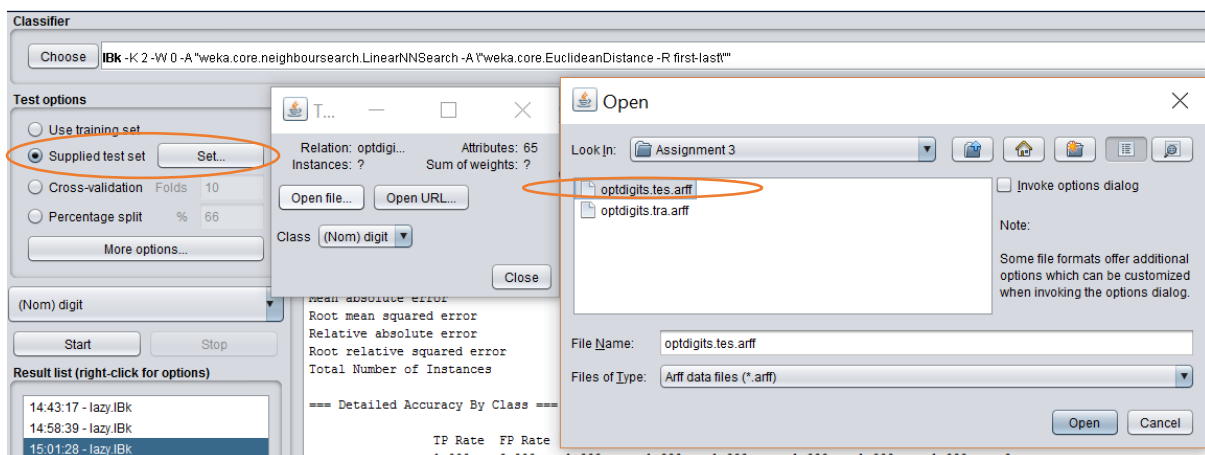
Step1: Weka Explorer -> Open file “optdigits.tra.arff”.

Step2: Classify -> Choose lazy -> IBk classifier.

Step3: Set kNN=1 -> Choose “EuclideanDistance”.



Step4: Select supplied test set as “optdigits.tes.arff”.



Step5: Start classification.

Step6: Change the number of K of IBk classifier.

3.2 Results

I use 3823 training instances to build the model and evaluate it with 1797 test instances. I select IBk (k-nearest neighbor) classifier as the instance-based classifier.

1. k=1, using 1 nearest neighbor(s) for classification

```

=== Summary ===
Correctly Classified Instances      1760      97.941 %
Incorrectly Classified Instances     37      2.059 %
Kappa statistic                    0.9771
Mean absolute error                 0.0046
Root mean squared error             0.0641
Relative absolute error             2.5427 %
Root relative squared error         21.3641 %
Total Number of Instances          1797

=== Confusion Matrix ===
  a  b  c  d  e  f  g  h  i  j  <-- classified as
178  0  0  0  0  0  0  0  0  0  | a = 0
  0 181  0  0  0  0  0  0  1  0  | b = 1
  0  2 175  0  0  0  0  0  0  0  | c = 2
  0  0  0 179  0  0  0  2  0  2  | d = 3
  0  2  0  0 178  0  0  0  1  0  | e = 4
  0  0  0  0  1 179  0  0  0  2  | f = 5
  0  0  0  0  0  0 181  0  0  0  | g = 6
  0  0  0  0  0  0  0 176  0  3  | h = 7
  0  9  0  1  0  0  0  0 163  1  | i = 8
  0  0  0  3  2  2  0  0  3 170  | j = 9

```

The actual accuracy rate is 97.941%, and the given accuracy rate is 98.00%. They are much similar. Based on values of **F-Measure**, it can be found that digit “8” and “9” are more difficult to distinguish, and digit “0” and “6” are easier to distinguish than others.

2. k=2, using 2 nearest neighbor(s) for classification

```

=== Summary ===
Correctly Classified Instances      1750      97.3845 %
Incorrectly Classified Instances     47      2.6155 %
Kappa statistic                    0.9709
Mean absolute error                 0.0052
Root mean squared error             0.0597
Relative absolute error             2.8784 %
Root relative squared error         19.8852 %
Total Number of Instances          1797

=== Confusion Matrix ===
  a  b  c  d  e  f  g  h  i  j  <-- classified as
178  0  0  0  0  0  0  0  0  0  | a = 0
  0 182  0  0  0  0  0  0  0  0  | b = 1
  0  4 173  0  0  0  0  0  0  0  | c = 2
  0  0  0 181  0  0  0  2  0  0  | d = 3
  0  2  0  0 178  0  0  0  1  0  | e = 4
  0  0  0  1  1 179  0  0  0  1  | f = 5
  1  0  0  0  0  1 179  0  0  0  | g = 6
  0  0  0  0  1  0  0 176  1  1  | h = 7
  0 10  1  3  0  0  0  0 159  1  | i = 8
  0  0  0  3  3  6  0  0  3 165  | j = 9

```

The actual accuracy rate is 97.3845%, and the given accuracy rate is 97.38%. They are much similar. Based on values of **F-Measure**, it can be found that digit “8” and “9” are more difficult to distinguish, and digit “0” and “6” are easier to distinguish than others.

3. k=3, using 3 nearest neighbor(s) for classification

```

=== Summary ===
Correctly Classified Instances      1758      97.8297 %
Incorrectly Classified Instances     39      2.1703 %
Kappa statistic                    0.9759
Mean absolute error                 0.0058
Root mean squared error             0.0599
Relative absolute error             3.2429 %
Root relative squared error         19.9604 %
Total Number of Instances          1797

=== Confusion Matrix ===
  a  b  c  d  e  f  g  h  i  j  <-- classified as
178  0  0  0  0  0  0  0  0  0  | a = 0
  0 180  0  0  0  0  1  0  1  0  | b = 1
  0  4 173  0  0  0  0  0  0  0  | c = 2
  0  0  0 181  0  0  0  1  1  0  | d = 3
  0  2  0  0 178  0  0  0  1  0  | e = 4
  0  0  0  1  1 179  0  0  0  1  | f = 5
  0  0  0  0  0  0 181  0  0  0  | g = 6
  0  0  0  0  0  0  0 172  1  6  | h = 7
  0  9  0  1  0  0  0  0 162  2  | i = 8
  0  0  0  4  0  1  0  0  1 174  | j = 9

```

The actual accuracy rate is 97.8297%, and the given accuracy rate is 97.83%. They are much similar. Based on values of **F-Measure**, it can be found that digit “8” and “1” are more difficult to distinguish, and digit “0” and “6” are easier to distinguish than others.

4. k=4, using 4 nearest neighbor(s) for classification

```

==== Confusion Matrix ====

=== Summary ===
Correctly Classified Instances      1757      97.7741 %
Incorrectly Classified Instances     40      2.2259 %
Kappa statistic                     0.9753
Mean absolute error                  0.0064
Root mean squared error              0.0592
Relative absolute error              3.5658 %
Root relative squared error          19.7464 %
Total Number of Instances           1797

      a  b  c  d  e  f  g  h  i  j  <-- classified as
178  0  0  0  0  0  0  0  0  0  | a = 0
0 182  0  0  0  0  0  0  0  0  | b = 1
0  4 173  0  0  0  0  0  0  0  | c = 2
0  1  0 179  0  0  0  2  1  0  | d = 3
0  2  0  0 179  0  0  0  0  0  | e = 4
0  0  0  0  1 179  0  0  0  2  | f = 5
1  0  0  0  0  1 179  0  0  0  | g = 6
0  0  0  0  0  0  0 174  1  4  | h = 7
0 10  0  1  0  0  0  0 162  1  | i = 8
0  1  0  2  1  2  0  0  2 172  | j = 9

```

The actual accuracy rate is 97.7741%, and the given accuracy rate is 97.61%. They are much similar. Based on values of **F-Measure**, it can be found that digit “8” and “1” are more difficult to distinguish, and digit “0” and “6” are easier to distinguish than others.

5. k=5, using 5 nearest neighbor(s) for classification

```

==== Confusion Matrix ====

=== Summary ===
Correctly Classified Instances      1760      97.941 %
Incorrectly Classified Instances     37      2.059 %
Kappa statistic                     0.9771
Mean absolute error                  0.0065
Root mean squared error              0.058
Relative absolute error              3.6139 %
Root relative squared error          19.3273 %
Total Number of Instances           1797

      a  b  c  d  e  f  g  h  i  j  <-- classified as
178  0  0  0  0  0  0  0  0  0  | a = 0
0 182  0  0  0  0  0  0  0  0  | b = 1
0  3 174  0  0  0  0  0  0  0  | c = 2
0  1  0 177  0  1  0  1  3  0  | d = 3
0  1  0  0 179  0  0  0  1  0  | e = 4
0  0  0  0  1 180  0  0  0  1  | f = 5
0  0  0  0  0  0 181  0  0  0  | g = 6
0  0  0  0  0  0  0 172  1  6  | h = 7
0  8  0  1  0  1  0  0 163  1  | i = 8
0  0  0  2  1  1  0  0  2 174  | j = 9

```

The actual accuracy rate is 97.941%, and the given accuracy rate is 97.89%. They are much similar. Based on values of **F-Measure**, it can be found that digit “8” and “9” are more difficult to distinguish, and digit “0” and “6” are easier to distinguish than others.

6. k=6, using 6 nearest neighbor(s) for classification

```

==== Confusion Matrix ====

=== Summary ===
Correctly Classified Instances      1755      97.6628 %
Incorrectly Classified Instances     42      2.3372 %
Kappa statistic                     0.974
Mean absolute error                  0.007
Root mean squared error              0.0599
Relative absolute error              3.8725 %
Root relative squared error          19.9522 %
Total Number of Instances           1797

      a  b  c  d  e  f  g  h  i  j  <-- classified as
178  0  0  0  0  0  0  0  0  0  | a = 0
0 182  0  0  0  0  0  0  0  0  | b = 1
0  4 173  0  0  0  0  0  0  0  | c = 2
0  1  0 177  0  1  0  2  2  0  | d = 3
0  2  0  0 179  0  0  0  0  0  | e = 4
0  0  0  0  1 179  0  0  0  2  | f = 5
0  0  0  0  0  1 180  0  0  0  | g = 6
0  0  0  0  0  0  0 175  1  3  | h = 7
0 10  0  2  0  0  0  0 161  1  | i = 8
0  3  0  2  0  2  0  0  2 171  | j = 9

```

The actual accuracy rate is 97.6628%, and the given accuracy rate is 97.77%. They are much similar. Based on values of **F-Measure**, it can be found that digit “8” and “1” are more difficult to distinguish, and digit “0” and “6” are easier to distinguish than others.

7. k=7, using 7 nearest neighbor(s) for classification

```

==== Confusion Matrix ====

=== Summary ===
Correctly Classified Instances      1753      97.5515 %
Incorrectly Classified Instances     44      2.4485 %
Kappa statistic                     0.9728
Mean absolute error                  0.0075
Root mean squared error              0.0613
Relative absolute error              4.1874 %
Root relative squared error          20.4433 %
Total Number of Instances           1797

      a  b  c  d  e  f  g  h  i  j  <-- classified as
178  0  0  0  0  0  0  0  0  0  | a = 0
0 180  0  0  0  0  1  0  1  0  | b = 1
0  4 173  0  0  0  0  0  0  0  | c = 2
0  1  0 177  0  1  0  2  2  0  | d = 3
0  2  0  0 179  0  0  0  0  0  | e = 4
0  0  0  0  1 179  0  0  0  2  | f = 5
0  0  0  0  0  1 180  0  0  0  | g = 6
0  0  0  0  0  0  0 174  2  3  | h = 7
0 10  0  2  0  0  0  0 160  2  | i = 8
0  1  0  2  0  1  0  0  3 173  | j = 9

```

The actual accuracy rate is 97.5515%, and the given accuracy rate is 97.66%. They are much similar. Based on values of **F-Measure**, it can be found that digit “8” and “1” are more difficult to distinguish, and digit “0” and “6” are easier to distinguish than others.

8. k=8, using 8 nearest neighbor(s) for classification

```

==== Confusion Matrix ====

a  b  c  d  e  f  g  h  i  j  <-- classified as
178 0  0  0  0  0  0  0  0  0 | a = 0
0 181 0  0  0  0  1  0  0  0 | b = 1
0  4 172 0  0  0  0  1  0  0 | c = 2
0  1  0 178 0  1  0  2  1  0 | d = 3
0  2  0  0 179 0  0  0  0  0 | e = 4
0  0  0  0  1 179 0  0  0  2 | f = 5
0  0  0  0  0  1 180 0  0  0 | g = 6
0  0  0  0  0  0  0 175  1  3 | h = 7
0 11  0  2  0  0  0  0 159  2 | i = 8
0  2  0  3  0  1  0  0  2 172 | j = 9

=== Summary ===
Correctly Classified Instances      1753      97.5515 %
Incorrectly Classified Instances     44      2.4485 %
Kappa statistic                    0.9728
Mean absolute error                 0.008
Root mean squared error             0.0616
Relative absolute error             4.4436 %
Root relative squared error        20.5372 %
Total Number of Instances          1797

```

The actual accuracy rate is 97.5515%, and the given accuracy rate is 97.66%. They are much similar. Based on values of **F-Measure**, it can be found that digit “8” and “1” are more difficult to distinguish, and digit “0” and “6” are easier to distinguish than others.

9. k=9, using 9 nearest neighbor(s) for classification

```

==== Confusion Matrix ====

a  b  c  d  e  f  g  h  i  j  <-- classified as
178 0  0  0  0  0  0  0  0  0 | a = 0
0 180 0  0  0  0  1  0  1  0 | b = 1
0  4 172 0  0  0  0  1  0  0 | c = 2
0  1  0 179 0  1  0  1  1  0 | d = 3
0  2  0  0 179 0  0  0  0  0 | e = 4
0  0  0  0  1 179 0  0  0  2 | f = 5
0  0  0  0  0  0 181 0  0  0 | g = 6
0  0  0  0  0  0  0 175  1  3 | h = 7
0  9  0  2  0  0  0  0 161  2 | i = 8
0  0  0  2  0  1  0  0  2 175 | j = 9

=== Summary ===
Correctly Classified Instances      1759      97.8854 %
Incorrectly Classified Instances     38      2.1146 %
Kappa statistic                    0.9765
Mean absolute error                 0.0082
Root mean squared error             0.0614
Relative absolute error             4.5531 %
Root relative squared error        20.483 %
Total Number of Instances          1797

```

The actual accuracy rate is 97.8854%, and the given accuracy rate is 97.72%. They are much similar. Based on values of **F-Measure**, it can be found that digit “8” and “1” are more difficult to distinguish, and digit “0” and “6” are easier to distinguish than others.

10. k=10, using 10 nearest neighbor(s) for classification

```

==== Confusion Matrix ====

a  b  c  d  e  f  g  h  i  j  <-- classified as
178 0  0  0  0  0  0  0  0  0 | a = 0
0 181 0  0  0  0  1  0  0  0 | b = 1
0  4 172 0  0  0  0  1  0  0 | c = 2
0  1  0 179 0  0  0  1  2  0 | d = 3
0  2  0  0 179 0  0  0  0  0 | e = 4
0  0  0  0  1 179 0  0  0  2 | f = 5
0  0  0  0  0  1 180 0  0  0 | g = 6
0  0  0  0  0  0  0 175  1  3 | h = 7
0 10  0  2  0  0  0  0 160  2 | i = 8
0  1  0  3  0  3  0  0  3 170 | j = 9

=== Summary ===
Correctly Classified Instances      1753      97.5515 %
Incorrectly Classified Instances     44      2.4485 %
Kappa statistic                    0.9728
Mean absolute error                 0.0085
Root mean squared error             0.0615
Relative absolute error             4.7275 %
Root relative squared error        20.5054 %
Total Number of Instances          1797

```

The actual accuracy rate is 97.5515%, and the given accuracy rate is 97.55%. They are much similar. Based on values of **F-Measure**, it can be found that digit “8” and “1” are more difficult to distinguish, and digit “0” and “6” are easier to distinguish than others.

11. k=11, using 11 nearest neighbor(s) for classification

```

=== Confusion Matrix ===

=== Summary ===
Correctly Classified Instances      1757      97.7741 %
Incorrectly Classified Instances      40      2.2259 %
Kappa statistic                    0.9753
Mean absolute error                  0.009
Root mean squared error              0.0628
Relative absolute error              4.9967 %
Root relative squared error          20.9356 %
Total Number of Instances           1797

      a  b  c  d  e  f  g  h  i  j  <-- classified as
178  0  0  0  0  0  0  0  0  0  | a = 0
  0 181  0  0  0  0  1  0  0  0  | b = 1
  0  4 172  0  0  0  0  1  0  0  | c = 2
  0  1  0  0 180  0  0  0  0  0  | d = 3
  0  0  0  0  1 179  0  0  0  2  | e = 4
  0  0  0  0  0  1 180  0  0  0  | g = 6
  0  0  0  0  0  0  0 175  1  3  | h = 7
  0  9  0  2  0  0  0  0 161  2  | i = 8
  0  0  0  3  0  2  0  0  2 173  | j = 9

```

The actual accuracy rate is 97.7741%, and the given accuracy rate is 97.89%. They are much similar. Based on values of **F-Measure**, it can be found that digit “8” and “1” are more difficult to distinguish, and digit “0” and “6” are easier to distinguish than others.

In general, the following table is the comparison of actual and given accuracy rates. They are similar.

	Actual accuracy rate	Given accuracy rate in the original file
k=1	97.941%	98.00%
k=2	97.3845%	97.38%
k=3	97.8297%	97.83%
k=4	97.7741%	97.61%
k=5	97.941%	97.89%
k=6	97.6628%	97.77%
k=7	97.5515%	97.66%
k=8	97.5515%	97.66%
k=9	97.8854%	97.72%
k=10	97.5515%	97.55%
k=11	97.7741%	97.89%

For KNN method, based on values of F-Measure, it can be found that digit “8”, “1”, and “9” are more difficult to distinguish, and digit “0” and “6” are easier to distinguish than others.

3.3 Identify a Misclassified Instance

For IBk, k=5, the confusion matrix is shown below. I pick one instance that represents digit “8” but is misclassified as digit “9”, circled below.

```


=== Confusion Matrix ===

      a  b  c  d  e  f  g  h  i  j  <-- classified as
178  0  0  0  0  0  0  0  0  0  | a = 0
  0 182  0  0  0  0  0  0  0  0  | b = 1
  0  3 174  0  0  0  0  0  0  0  | c = 2
  0  1  0 177  0  1  0  1  3  0  | d = 3
  0  1  0  0 179  0  0  0  1  0  | e = 4
  0  0  0  0  1 180  0  0  0  1  | f = 5
  0  0  0  0  0  0 181  0  0  0  | g = 6
  0  0  0  0  0  0  0 172  1  6  | h = 7
  0  8  0  1  0  1  0  0 163  ①  | i = 8
  0  0  0  2  1  1  0  0  2 174  | j = 9

```


Right click on the model, and choose “Visualize classifier errors”.

Find the green square box which shows “predicted digit : 9; digit : 8”. The detailed information is shown below. It is instance No. 1469.



Modify the Python program to display the instance No. 1469.

```
#check out instance number 1469
# it was predicted 9 but actually an 8 (based on IBk, k=5)
lines = lines[1468:1469]
```



It can be seen that the digit “8” of instance No. 1469 indeed looks like digit “9”.

4 Other Classifiers Comparison

1. Function classifier: Logistic Regression Model (Non-instance based)

Accuracy rate: 92.1536%.

```
=== Summary ===
Correctly Classified Instances      1656      92.1536 %
Incorrectly Classified Instances    141       7.8464 %
Kappa statistic                    0.9128
Mean absolute error                 0.0157
Root mean squared error             0.1236
Relative absolute error             8.6954 %
Root relative squared error         41.2095 %
Total Number of Instances          1797

=== Confusion Matrix ===
      a  b  c  d  e  f  g  h  i  j  <-- classified as
170  0  1  0  1  3  1  1  0  1 | a = 0
 3 169  0  0  4  0  1  1  4  0 | b = 1
 0  2 169  1  0  1  1  0  3  0 | c = 2
 0  0  3 167  0  5  0  1  5  2 | d = 3
 1  3  0  0 168  0  0  3  1  5 | e = 4
 0  0  2  5  0 171  0  2  1  1 | f = 5
 0  0  1  0  8  0 171  0  1  0 | g = 6
 3  0  0  0  1  5  0 165  2  3 | h = 7
 0  6  3  7  1  2  3  2 146  4 | i = 8
 6  0  0  6  0  1  1  3  3 160 | j = 9
```

2. Trees classifier: J48 (Non-instance based)

Accuracy rate: 85.754%.


```

=== Summary ===
Correctly Classified Instances      1541      85.754 %
Incorrectly Classified Instances    256      14.246 %
Kappa statistic                    0.8417
Mean absolute error                0.0306
Root mean squared error            0.1614
Relative absolute error            16.9876 %
Root relative squared error        53.8008 %
Total Number of Instances         1797

=== Confusion Matrix ===
      a  b  c  d  e  f  g  h  i  j  <-- classified as
169  0  0  0  5  2  0  0  1  1 | a = 0
  0 164  1  2  0  0  2  8  2  3 | b = 1
  1  4 149  6  4  0  5  0  8  0 | c = 2
  0 11  6 147  1  0  0  4  7  7 | d = 3
  3 11  1  1 149  0  2  3 10  1 | e = 4
  2  1  1  4  5 159  4  1  1  4 | f = 5
  0  2  0  0  2  1 170  0  6  0 | g = 6
  0  0  0  1  8  0  1 141 13 15 | h = 7
  0  8 10  3  2  4  1  1 143  2 | i = 8
  3  1  7  5  5  5  0  3  1 150 | j = 9

```

3. Instance-based classifier: IBK (kNN, k=2)

Accuracy rate: 97.3845%

```

=== Summary ===
Correctly Classified Instances      1750      97.3845 %
Incorrectly Classified Instances    47      2.6155 %
Kappa statistic                    0.9709
Mean absolute error                0.0052
Root mean squared error            0.0597
Relative absolute error            2.8784 %
Root relative squared error        19.8852 %
Total Number of Instances         1797

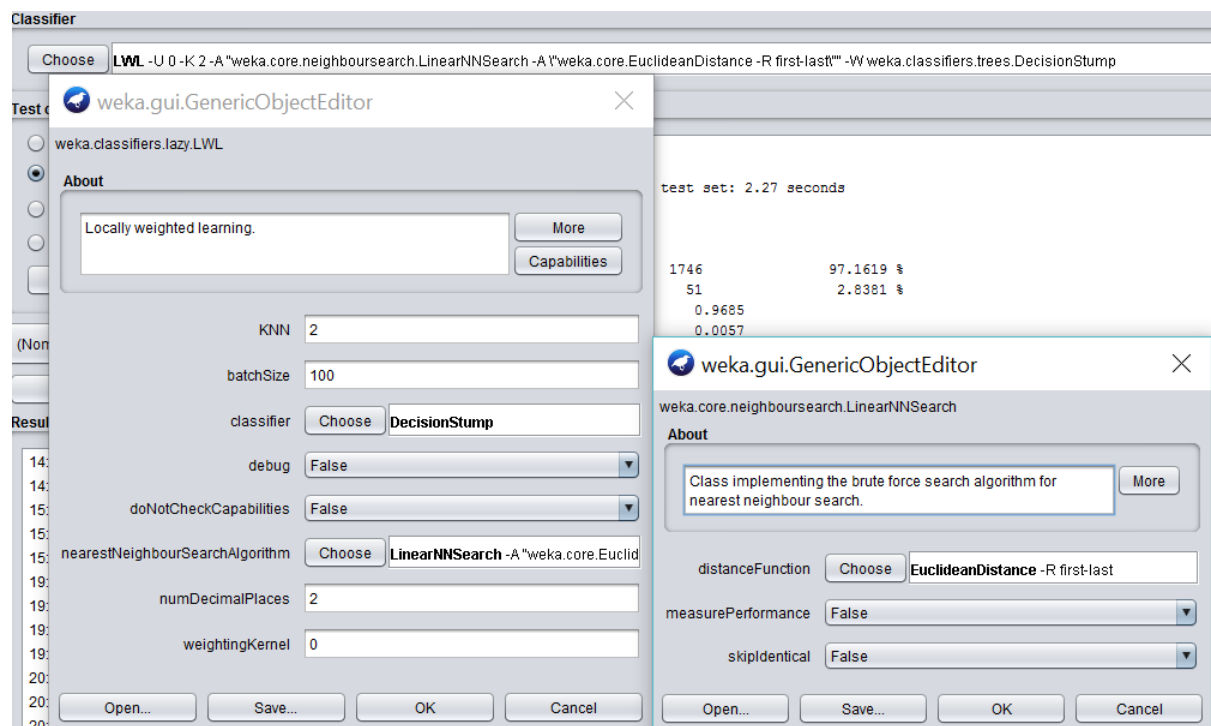
=== Confusion Matrix ===
      a  b  c  d  e  f  g  h  i  j  <-- classified as
178  0  0  0  0  0  0  0  0  0 | a = 0
  0 182  0  0  0  0  0  0  0  0 | b = 1
  0  4 173  0  0  0  0  0  0  0 | c = 2
  0  0  0 181  0  0  0  2  0  0 | d = 3
  0  2  0  0 178  0  0  0  1  0 | e = 4
  0  0  0  1  1 179  0  0  0  1 | f = 5
  1  0  0  0  0  1 179  0  0  0 | g = 6
  0  0  0  0  1  0  0 176  1  1 | h = 7
  0 10  1  3  0  0  0  0 159  1 | i = 8
  0  0  0  3  3  6  0  0  3 165 | j = 9

```

4. Instance-based classifier: LWL (k=2)

Locally weighted learning. Uses an instance-based algorithm to assign instance weights which are then used by a specified WeightedInstancesHandler.

Accuracy rate: 97.1619%



```

=== Summary ===
Correctly Classified Instances      1746      97.1619 %
Incorrectly Classified Instances     51      2.8381 %
Kappa statistic                    0.9685
Mean absolute error                 0.0057
Root mean squared error             0.0753
Relative absolute error             3.1534 %
Root relative squared error        25.1133 %
Total Number of Instances         1797

=== Confusion Matrix ===
      a  b  c  d  e  f  g  h  i  j  <-- classified as
178    0  0  0  0  0  0  0  0  0  a = 0
  0 180    0  0  0  0  2  0  0  0  b = 1
  0   3 174    0  0  0  0  0  0  0  c = 2
  0   0  0 177    0  1  0  1  2  2  d = 3
  0   2   0   0 178    0  0  0  1  0  e = 4
  1   0   0   0   0  0 180    0  0  0  f = 5
  0   0   0   0   1   0   0 172    1   5  g = 6
  0   9   1   0   0   0   0   0 160    4  h = 7
  0   0   0   2   3   5   0   0   2 168  i = 8
                                j = 9

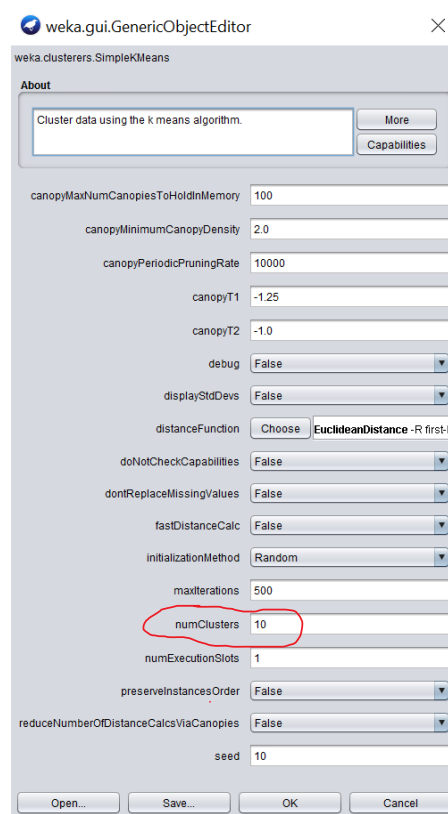
```

In general, instance-based classifiers perform better than non-instance-based classifiers in this problem domain. For instance-based classifiers, they do not conduct rule induction, instead, they compare instances in the test set with instances in the training test. They construct assumptions based on training set, and apply the assumptions to the test set. If there are many overlapping instances or patterns between the training and test sets, the accuracy rate will be high. In this digit recognition scenario, the thousands of instances come from 43 people (30 training set, 13 test set), so there exists some similarity between instances from the 13 people and instances from the previous 30 people. Therefore, the assumptions based on the training set could be well applied to the test set, and the corresponding accuracy rate could also be high.

For other non-instance based classifiers such as logistic regression and J48 (decision tree), they fit (parametric or non-parametric) models based on the training set and predict instances in the test set based on the models. Some statistical issues in these models may cause lower accuracy rates than those of instance-based classifiers.

5 Comparison of Classifiers and Clusterers

1. SimpleKMeans Cluster -> 10 clusters



- Result of Supplied test set

All the nine digits should be even distributed in the test set. However, the clustered result of the test set is quite not uniform distributed.

```
=== Evaluation on test set ===
Clustered Instances
```

```
0      190 ( 11%)
1      168 (  9%)
2      168 (  9%)
3      178 ( 10%)
4      181 ( 10%)
5      100 (  6%)
6      369 ( 21%)
7       72 (  4%)
8      255 ( 14%)
9      116 (  6%)
```

- Result of training set cluster evaluation

Almost one third of instances are incorrectly clustered. The accuracy rate is about 70%.

```
=== Model and evaluation on training set ===
```

```
Clustered Instances
```

```
0      456 ( 12%)
1      316 (  8%)
2      315 (  8%)
3      372 ( 10%)
4      391 ( 10%)
5      275 (  7%)
6      768 ( 20%)
7      161 (  4%)
8      537 ( 14%)
9      232 (  6%)
```

```
Class attribute: digit
```

```
Classes to Clusters:
```

```
0  1  2  3  4  5  6  7  8  9  <-- assigned to cluster
0  1  0  372  1  1  0  0  1  0 | 0
2  0  1  0  2  137  9  5  219  14 | 1
7  0  0  0  1  0  1  142  14  215 | 2
12  0  6  0  0  4  349  4  11  3 | 3
32 314  7  0  4  27  0  0  3  0 | 4
0  0  294  0  1  4  71  0  6  0 | 5
0  1  0  0  374  0  0  0  2  0 | 6
375 0  0  0  0  5  0  0  7  0 | 7
3  0  3  0  8  5  88  2  271  0 | 8
25 0  4  0  0  92  250  8  3  0 | 9
```

Cluster mode

☐ Use training set

☐ Supplied test set

☐ Percentage split %

☒ Classes to clusters evaluation

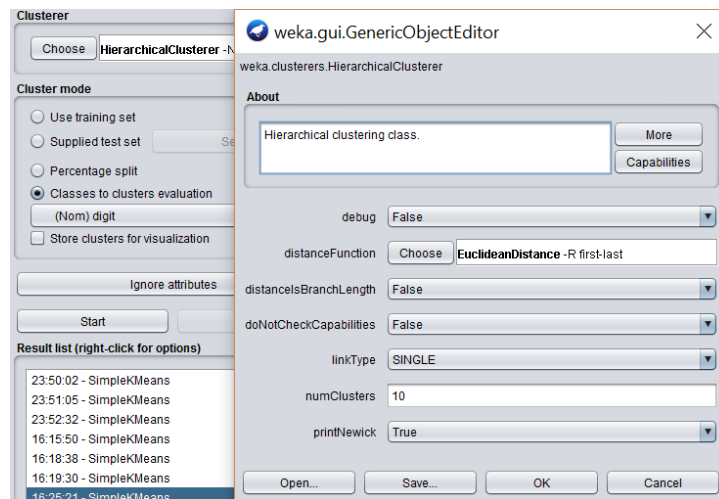
▼

☐ Store clusters for visualization

```
Cluster 0 <-- 7
Cluster 1 <-- 4
Cluster 2 <-- 5
Cluster 3 <-- 0
Cluster 4 <-- 6
Cluster 5 <-- 1
Cluster 6 <-- 3
Cluster 7 <-- 9
Cluster 8 <-- 8
Cluster 9 <-- 2
```

```
Incorrectly clustered instances :      1114.0    29.1394 %
```

2. Hierarchical Cluster -> 10 clusters



- Result of training set cluster evaluation (change linkType to “Average”)
In the training set, only 8 digit numbers have corresponding classes. 38.8438% instances are incorrectly clustered. The accuracy rate is about 61%.

```

=== Model and evaluation on training set ===

Clustered Instances

0      753 ( 20%)
1      388 ( 10%)
2      326 (  9%)
3      395 ( 10%)
4      994 ( 26%)
5      601 ( 16%)
6      327 (  9%)
7       16 (  0%)
8       20 (  1%)
9        3 (  0%)

Class attribute: digit
Classes to Clusters:

  0  1  2  3  4  5  6  7  8  9 <-- assigned to cluster
374 0 0 0 0 0 1 1 0 0 0 0
  0 0 155 25 14 195 0 0 0 0 1
  1 11 0 356 4 8 0 0 0 0 2
  0 0 1 1 368 19 0 0 0 0 3
  3 1 38 0 0 0 325 0 20 0 4
  0 0 2 0 357 1 0 16 0 0 5
375 0 1 0 0 1 0 0 0 0 0 6
  0 369 14 0 0 0 1 0 0 3 7
  0 0 4 0 2 374 0 0 0 0 8
  0 7 111 13 249 2 0 0 0 0 9

Cluster 0 <-- 6
Cluster 1 <-- 7
Cluster 2 <-- 1
Cluster 3 <-- 2
Cluster 4 <-- 3
Cluster 5 <-- 8
Cluster 6 <-- 4
Cluster 7 <-- 5
Cluster 8 <-- No class
Cluster 9 <-- No class

Incorrectly clustered instances :      1485.0   38.8438 %

```

The performance of the clustering is not quite surprising.

For classification, there exists gold standard as the class labels for supervised learning.

For clustering, it is unsupervised learning without any gold standard. Similar instances are clustered as one group, and the class label is ignored. It should be less accurate than classification since similar instances probably may come from the same person but not always the same digit number. The personal handwritten style may affect the cluster results. Clustering may be based on the characteristics of both digit numbers and different people handwritten styles in this domain.

6 10-fold Cross-Validation of Training set

I pick the best classifier “IBK, K=5” so far for evaluating using 10-fold cross-validation of the training set. The accuracy rate is 98.5613%, higher than the previous provided test set (97.941%).

```

=== Stratified cross-validation ===
=== Summary ===

Correctly Classified Instances      3768           98.5613 %
Incorrectly Classified Instances     55           1.4387 %
Kappa statistic                     0.984
Mean absolute error                  0.005
Root mean squared error              0.0501
Relative absolute error              2.7818 %
Root relative squared error          16.6929 %
Total Number of Instances           3823

```

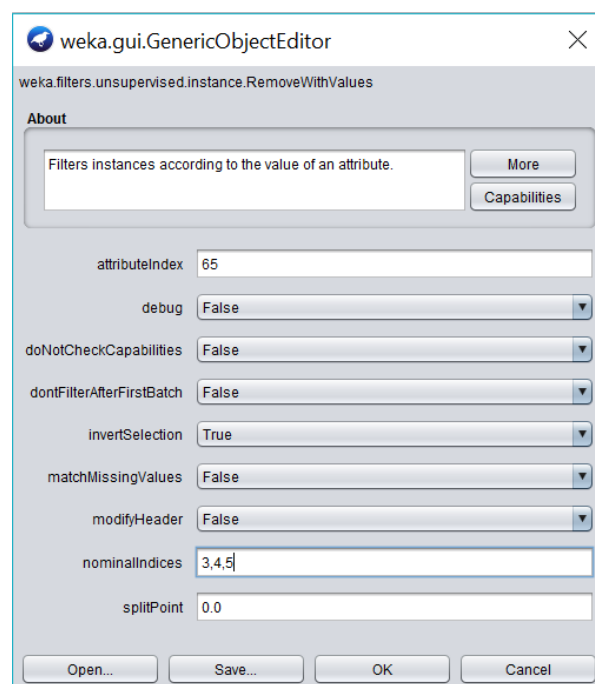
For the 10-fold cross-validation, each subset in the 10 subsets of the training set has been used as test set, and then 10 models have been evaluated on the 10 subsets. The classified accuracy rate is an average of accuracy rates of the 10 models. Compared to evaluating model on another test set, 10-fold cross-validation of the training set can decrease the bias between actual training/test set between the whole dataset, since each subset has been used as training or test set in the 10 models, like sampling method in statistics. Therefore, 10-fold cross-validation of the training set performs higher accuracy rate. It can also reduce overfitting to some extent.

7 Accuracy VS. Number of Classes

I suppose the question is to find a relationship between number of class labels and the accuracy rate of classification. Expanding the dataset means adding more instances with more class labels to it.

7.1 Methods

Step1: “optdigits.tra.arff” -> Preprocess -> RemoveWithValues -> select instances of certain digit numbers



Step2: Classify -> IBK (k=5) -> 10-fold cross-validation

7.2 Results

1. 3 classes (digits “2”, “3”, “4”); Accuracy rate: 99.827%

No.	Label	Count	Weight	=== Summary ===		
1	0	0	0.0	Correctly Classified Instances	1154	99.827 %
2	1	0	0.0	Incorrectly Classified Instances	2	0.173 %
3	2	380	380.0	Kappa statistic	0.9974	
4	3	389	389.0	Mean absolute error	0.001	
5	4	387	387.0	Root mean squared error	0.0158	
6	5	0	0.0	Relative absolute error	0.7229 %	
7	6	0	0.0	Root relative squared error	6.1201 %	
8	7	0	0.0	Total Number of Instances	1156	
9	8	0	0.0			
10	9	0	0.0			

2. 4 classes (digits “2”, “3”, “4”, “5”); Accuracy rate: 99.6736%

No.	Label	Count	Weight	=== Summary ===		
1	0	0	0.0	Correctly Classified Instances	1527	99.6736 %
2	1	0	0.0	Incorrectly Classified Instances	5	0.3264 %
3	2	380	380.0	Kappa statistic	0.9956	
4	3	389	389.0	Mean absolute error	0.0014	
5	4	387	387.0	Root mean squared error	0.0234	
6	5	376	376.0	Relative absolute error	0.9546 %	
7	6	0	0.0	Root relative squared error	8.5528 %	
8	7	0	0.0	Total Number of Instances	1532	
9	8	0	0.0			
10	9	0	0.0			

3. 5 classes (digits “2”, “3”, “4”, “5”, “6”); Accuracy rate: 99.5809%

No.	Label	Count	Weight	=== Summary ===		
1	0	0	0.0	Correctly Classified Instances	1901	99.5809 %
2	1	0	0.0	Incorrectly Classified Instances	8	0.4191 %
3	2	380	380.0	Kappa statistic	0.9948	
4	3	389	389.0	Mean absolute error	0.0014	
5	4	387	387.0	Root mean squared error	0.0253	
6	5	376	376.0	Relative absolute error	0.8889 %	
7	6	377	377.0	Root relative squared error	8.9529 %	
8	7	0	0.0	Total Number of Instances	1909	
9	8	0	0.0			
10	9	0	0.0			

4. 6 classes (digits “1”, “2”, “3”, “4”, “5”, “6”); Accuracy rate: 99.6084%

No.	Label	Count	Weight	=== Summary ===		
1	0	0	0.0	Correctly Classified Instances	2289	99.6084 %
2	1	389	389.0	Incorrectly Classified Instances	9	0.3916 %
3	2	380	380.0	Kappa statistic	0.9953	
4	3	389	389.0	Mean absolute error	0.0016	
5	4	387	387.0	Root mean squared error	0.0274	
6	5	376	376.0	Relative absolute error	0.9699 %	
7	6	377	377.0	Root relative squared error	9.4776 %	
8	7	0	0.0	Total Number of Instances	2298	
9	8	0	0.0			
10	9	0	0.0			

5. 7 classes (digits “0”, “1”, “2”, “3”, “4”, “5”, “6”); Accuracy rate: 99.4764%

No.	Label	Count	Weight	=== Summary ===		
1	0	376	376.0	Correctly Classified Instances	2660	99.4764 %
2	1	389	389.0	Incorrectly Classified Instances	14	0.5236 %
3	2	380	380.0	Kappa statistic	0.9939	
4	3	389	389.0	Mean absolute error	0.0018	
5	4	387	387.0	Root mean squared error	0.0311	
6	5	376	376.0	Relative absolute error	1.0649 %	
7	6	377	377.0	Root relative squared error	10.6205 %	
8	7	0	0.0	Total Number of Instances	2674	
9	8	0	0.0			
10	9	0	0.0			

6. 8 classes (digits “0”, “1”, “2”, “3”, “4”, “5”, “6”, “8”); Accuracy rate: 99.1159%

=== Summary ===			
No.	Label	Count	Weight
1	0	376	376.0
2	1	389	389.0
3	2	380	380.0
4	3	389	389.0
5	4	387	387.0
6	5	376	376.0
7	6	377	377.0
8	7	0	0.0
9	8	380	380.0
10	9	0	0.0

Correctly Classified Instances	3027	99.1159 %
Incorrectly Classified Instances	27	0.8841 %
Kappa statistic	0.9899	
Mean absolute error	0.0033	
Root mean squared error	0.041	
Relative absolute error	1.8671 %	
Root relative squared error	13.8647 %	
Total Number of Instances	3054	

7. 9 classes (digits “0”, “1”, “2”, “3”, “4”, “5”, “6”, “7”, “8”); Accuracy rate: 98.9538%

=== Summary ===			
No.	Label	Count	Weight
1	0	376	376.0
2	1	389	389.0
3	2	380	380.0
4	3	389	389.0
5	4	387	387.0
6	5	376	376.0
7	6	377	377.0
8	7	387	387.0
9	8	380	380.0
10	9	0	0.0

Correctly Classified Instances	3405	98.9538 %
Incorrectly Classified Instances	36	1.0462 %
Kappa statistic	0.9882	
Mean absolute error	0.0032	
Root mean squared error	0.0415	
Relative absolute error	1.7784 %	
Root relative squared error	13.9148 %	
Total Number of Instances	3441	

8. 10 classes (digits “0”, “1”, “2”, “3”, “4”, “5”, “6”, “7”, “8”, “9”); Accuracy rate: 98.5613%

=== Summary ===			
No.	Label	Count	Weight
1	0	376	376.0
2	1	389	389.0
3	2	380	380.0
4	3	389	389.0
5	4	387	387.0
6	5	376	376.0
7	6	377	377.0
8	7	387	387.0
9	8	380	380.0
10	9	382	382.0

Correctly Classified Instances	3768	98.5613 %
Incorrectly Classified Instances	55	1.4387 %
Kappa statistic	0.984	
Mean absolute error	0.005	
Root mean squared error	0.0501	
Relative absolute error	2.7818 %	
Root relative squared error	16.6929 %	
Total Number of Instances	3823	

In general, a linear regression model is established based on the 8 pairs of values. (Much uncertainty)

Accuracy = 100.4319% - 0.1665* # of classes (not accurate)

```
lm(formula = accuracy ~ label, data = rate)
```

Residuals:


Min	1Q	Median	3Q	Max
-0.205617	-0.095584	-0.001263	0.059160	0.209987

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	100.43190	0.16117	623.13	1.15e-15 ***
label	-0.16650	0.02339	-7.12	0.000386 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.1516 on 6 degrees of freedom
Multiple R-squared: 0.8942, Adjusted R-squared: 0.8765
F-statistic: 50.69 on 1 and 6 DF, p-value: 0.0003862



If lower-case English alphabetical characters are included in the dataset, there will be 36 classes. Then the accuracy rate may be around 94.4% for IBK ($k=5$) 10-fold cross-validation on the training set, reduced by nearly 4% compared to the original 10 classes.

If all ASCII characters are included in the dataset, there will be 128 classes. Then the accuracy rate may be around 79% for IBK ($k=5$) 10-fold cross-validation on the training set, reduced by nearly 20% compared to the original 10 classes.

If all UTF-8 characters are included in the dataset, there will be 256 classes. Then the accuracy rate may be around 58% for IBK ($k=5$) 10-fold cross-validation on the training set, reduced by nearly 41% compared to the original 10 classes.

I just assume the accuracy rate is linear correlated to the number of classes. So, there could be high errors based on my prediction model. Generally, as the number of classes increases, the accuracy rate declines.

8 Pre-Processing Steps

I suppose the handwritten note only has digits on it. After scanning it into a pixel-based image, it may be possible to firstly recognize the image region for each digit number, and to divide the whole image into several sub-images which include only one digit for each image (removing some extra white space if necessary). Since there could be rotations for some sub-images, multiple tries could be done to capture better results. Besides, I suppose initially it is not compulsory to restrict the size of all sub-images to be the same. Maybe other methods can be utilized to modify these images in equal size in the following step. Furthermore, it may be possible to decrease noise through some programs.

Then for each sub-image with single digit, preprocessing programs available by NIST could be used to extract normalized bitmaps (black-and-white), as the paper stated. Then 32x32 bitmaps are divided into nonoverlapping blocks of 4x4 and the number of on pixels are counted in each block. This generates an input matrix of 8x8 where each element is an integer in the range 0 through 16. This reduces dimensionality and gives invariance to small distortions. The 64 values for each digit image should be recorded in order as one row (instance) in the dataset.

If supervised learning is needed, the class label (real digit number) of each instance (image) should also be recorded as the class attribute column in the dataset. Then a trained classifier can be applied to the dataset.